

# Assignment 4 MapReduce

*Task 1: Write a Map-Reduce Program to filter invalid record.*

Text file

```
[acadgild@localhost mapreduce]$ cat television.txt
Samsung|Optima|14|Madhya Pradesh|132401|14200
Onida|Lucid|18|Uttar Pradesh|232401|16200
Akai|Decent|16|Kerala|922401|12200
Lava|Attention|20|Assam|454601|24200
Zen|Super|14|Maharashtra|619082|9200
Samsung|Optima|14|Madhya Pradesh|132401|14200
Onida|Lucid|18|Uttar Pradesh|232401|16200
Onida|Decent|14|Uttar Pradesh|232401|16200
Onida|NA|16|Kerala|922401|12200
Lava|Attention|20|Assam|454601|24200
Zen|Super|14|Maharashtra|619082|9200
Samsung|Optima|14|Madhya Pradesh|132401|14200
NA|Lucid|18|Uttar Pradesh|232401|16200
Samsung|Decent|16|Kerala|922401|12200
Lava|Attention|20|Assam|454601|24200
Samsung|Super|14|Maharashtra|619082|9200
Samsung|Super|14|Maharashtra|619082|9200
Samsung|Super|14|Maharashtra|619082|9200[acadgild@localhost mapreduce]$
```

Command:

```
hadoop jar assign4.jar mrTask1.Task1 /mapreduce/television.txt /mapreduce/task1output
```

Code Snippet:

```
[acadgild@localhost mapreduce]$ hadoop jar assign4.jar mrTask1.Task1 /mapreduce/television.txt /mapreduce/task1output
19/01/11 06:09:14 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
19/01/11 06:09:15 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
19/01/11 06:09:17 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with
ToolRunner to remedy this.
19/01/11 06:09:17 INFO input.FileInputFormat: Total input paths to process : 1
19/01/11 06:09:17 INFO mapreduce.JobSubmitter: number of splits:1
19/01/11 06:09:18 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1547166974239_0001
19/01/11 06:09:18 INFO impl.YarnClientImpl: Submitted application application_1547166974239_0001
19/01/11 06:09:18 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1547166974239_0001/
19/01/11 06:09:18 INFO mapreduce.Job: Running job: job_1547166974239_0001
19/01/11 06:09:32 INFO mapreduce.Job: Job job_1547166974239_0001 running in uber mode : false
19/01/11 06:09:32 INFO mapreduce.Job: map 0% reduce 0%
19/01/11 06:09:39 INFO mapreduce.Job: map 100% reduce 0%
19/01/11 06:09:39 INFO mapreduce.Job: Job job_1547166974239_0001 completed successfully
19/01/11 06:09:39 INFO mapreduce.Job: Counters: 30
  File System Counters
    FILE: Number of bytes read=0
    FILE: Number of bytes written=107363
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=844
    HDFS: Number of bytes written=73
    HDFS: Number of read operations=5
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=4850
    Total time spent by all reducers in occupied slots (ms)=0
```

```

19/01/11 06:09:32 INFO mapreduce.Job: Job job_1547166974239_0001 running in uber mode : false
19/01/11 06:09:32 INFO mapreduce.Job: map 0% reduce 0%
19/01/11 06:09:39 INFO mapreduce.Job: map 100% reduce 0%
19/01/11 06:09:39 INFO mapreduce.Job: Job job_1547166974239_0001 completed successfully
19/01/11 06:09:39 INFO mapreduce.Job: Counters: 30
  File System Counters
    FILE: Number of bytes read=0
    FILE: Number of bytes written=107363
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=844
    HDFS: Number of bytes written=73
    HDFS: Number of read operations=5
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=4850
    Total time spent by all reduces in occupied slots (ms)=0
    Total time spent by all map tasks (ms)=4850
    Total vcore-milliseconds taken by all map tasks=4850
    Total megabyte-milliseconds taken by all map tasks=4966400
  Map-Reduce Framework
    Map input records=18
    Map output records=2
    Input split bytes=111
    Spilled Records=0
    Failed Shuffles=0
    Merged Map outputs=0
    GC time elapsed (ms)=66
    CPU time spent (ms)=510
    Physical memory (bytes) snapshot=106095712
    Virtual memory (bytes) snapshot=2056757248
    Total committed heap usage (bytes)=62900096
  File Input Format Counters
    Bytes Read=733
  File Output Format Counters
    Bytes Written=73
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost mapreduce]$ hadoop jar assign4.jar mrTask2.Task2 /mapreduce/television.txt /mapreduce/task2output
19/01/11 06:10:39 WARN util.NativeCodeLoader: Unable to load native hadoop library for your platform... using built-in code; performance will be degraded

```

Output Snippet:

```

[acadgild@localhost ~]$ hdfs dfs -cat /mapreduce/task2output
19/01/11 06:10:39 WARN util.NativeCodeLoader: Unable to load native hadoop library for your platform... using built-in code; performance will be degraded
Onida|NA|16|Kerala|922401|12200
NA|Lucid|18|Uttar Pradesh|232401|16200
[acadgild@localhost ~]$

```

## Task 2: Write a Map-Reduce Program to calculate the total unit sold for each company

Command:

```
hadoop jar assign4.jar mrTask2.Task2 /mapreduce/television.txt /mapreduce/task2output
```

Code Snippet:

```
[acadgild@localhost mapreduce]$ hadoop jar assign4.jar mrTask2.Task2 /mapreduce/television.txt /mapreduce/task2output
19/01/11 06:17:03 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
19/01/11 06:17:04 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
19/01/11 06:17:05 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with
ToolRunner to remedy this.
19/01/11 06:17:06 INFO input.FileInputFormat: Total input paths to process : 1
19/01/11 06:17:06 INFO mapreduce.JobSubmitter: number of splits:1
19/01/11 06:17:06 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1547166974239_0002
19/01/11 06:17:06 INFO impl.YarnClientImpl: Submitted application application_1547166974239_0002
19/01/11 06:17:07 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1547166974239_0002/
19/01/11 06:17:07 INFO mapreduce.Job: Running job: job_1547166974239_0002
19/01/11 06:17:16 INFO mapreduce.Job: Job job_1547166974239_0002 running in uber mode : false
19/01/11 06:17:22 INFO mapreduce.Job: map 0% reduce 0%
19/01/11 06:17:22 INFO mapreduce.Job: map 100% reduce 0%
19/01/11 06:17:32 INFO mapreduce.Job: map 100% reduce 100%
19/01/11 06:17:32 INFO mapreduce.Job: Job job_1547166974239_0002 completed successfully
19/01/11 06:17:33 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=225
    FILE: Number of bytes written=215827
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=844
    HDFS: Number of bytes written=43
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=4869
    Total time spent by all reduces in occupied slots (ms)=6575
    Total time spent by all map tasks (ms)=4869
    Total time spent by all reduce tasks (ms)=6575
    Total vcore-milliseconds taken by all map tasks=4869
    Total vcore-milliseconds taken by all reduce tasks=6575
    Total megabyte-milliseconds taken by all map tasks=4985856
    Total megabyte-milliseconds taken by all reduce tasks=6732800
  Map-Reduce Framework
    Map input records=18
    Total vcore-milliseconds taken by all reduce tasks=6575
    Total megabyte-milliseconds taken by all map tasks=4985856
    Total megabyte-milliseconds taken by all reduce tasks=6732800
  Map-Reduce Framework
    Map input records=18
    Map output records=18
    Map output bytes=183
    Map output materialized bytes=225
    Input split bytes=111
    Combine input records=0
    Combine output records=0
    Reduce input groups=0
    Reduce shuffle bytes=225
    Reduce input records=18
    Reduce output records=0
    Spilled Records=36
    Shuffled Maps =1
    Failed Shuffles=0
    Merged Map outputs=1
    GC time elapsed (ms)=196
    CPU time spent (ms)=1530
    Physical memory (bytes) snapshot=320380928
    Virtual memory (bytes) snapshot=4118196224
    Total committed heap usage (bytes)=222429184
  Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
  File Input Format Counters
    Bytes Read=733
  File Output Format Counters
    Bytes Written=43
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost mapreduce]$
```

Output Snippet:

```
[acadgild@localhost ~]$ hdfs dfs -cat /mapreduce/task2output/part-r-00000
19/01/11 06:20:08 WARN util.NativeCodeLoader: Unable to load native-hadoop
Akai      1
Lava      3
NA        1
Onida     4
Samsung   7
Zen       2
You have new mail in /var/spool/mail/acadgild
```

*Task 3: Write a Map-Reduce Program to calculate the total unit sold in each state for Onida company.*

Command:

```
hadoop jar assign4.jar mrTask3.Task3 /mapreduce/television.txt /mapreduce/task3output
```

Code Snippet:

```
[acadgild@localhost mapreduce]$ hadoop jar assign4.jar mrTask3.Task3 /mapreduce/television.txt /mapreduce/task3output
19/01/11 06:20:54 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
19/01/11 06:20:55 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
19/01/11 06:20:56 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application via
ToolRunner to remedy this.
19/01/11 06:20:57 INFO input.FileInputFormat: Total input paths to process : 1
19/01/11 06:20:57 INFO mapreduce.JobSubmitter: number of splits:1
19/01/11 06:20:57 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1547166974239_0003
19/01/11 06:20:57 INFO impl.YarnClientImpl: Submitted application application_1547166974239_0003
19/01/11 06:20:57 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1547166974239_0003/
19/01/11 06:20:57 INFO mapreduce.Job: Running job: job_1547166974239_0003
19/01/11 06:21:07 INFO mapreduce.Job: Job job_1547166974239_0003 running in uber mode : false
19/01/11 06:21:07 INFO mapreduce.Job:  map 0% reduce 0%
19/01/11 06:21:14 INFO mapreduce.Job:  map 100% reduce 0%
19/01/11 06:21:21 INFO mapreduce.Job:  map 100% reduce 100%
19/01/11 06:21:21 INFO mapreduce.Job: Job job_1547166974239_0003 completed successfully
19/01/11 06:21:21 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=387
    FILE: Number of bytes written=216137
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=844
    HDFS: Number of bytes written=64
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=4582
    Total time spent by all reduces in occupied slots (ms)=4945
    Total time spent by all map tasks (ms)=4582
    Total time spent by all reduce tasks (ms)=4945
    Total vcore-milliseonds taken by all map tasks=4582
    Total vcore-milliseonds taken by all reduce tasks=4945
    Total megabyte-milliseonds taken by all map tasks=4691968
    Total megabyte-milliseonds taken by all reduce tasks=5063680
  Map-Reduce Framework
    Map input records=18
    Map output records=22
```

```

    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=4582
    Total time spent by all reduces in occupied slots (ms)=4945
    Total time spent by all map tasks (ms)=4582
    Total time spent by all reduce tasks (ms)=4945
    Total vcore-milliseconds taken by all map tasks=4582
    Total vcore-milliseconds taken by all reduce tasks=4945
    Total megabyte-milliseconds taken by all map tasks=4691968
    Total megabyte-milliseconds taken by all reduce tasks=5063680
Map-Reduce Framework
  Map input records=18
  Map output records=22
  Map output bytes=337
  Map output materialized bytes=387
  Input split bytes=111
  Combine input records=0
  Combine output records=0
  Reduce input groups=5
  Reduce shuffle bytes=387
  Reduce input records=22
  Reduce output records=5
  Spilled Records=44
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=157
  CPU time spent (ms)=1320
  Physical memory (bytes) snapshot=320028672
  Virtual memory (bytes) snapshot=4118196224
  Total committed heap usage (bytes)=222429184
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=733
File Output Format Counters
  Bytes Written=64
[acadgild@localhost mapreduce]$

```

Output Snippet:

```

[acadgild@localhost ~]$ hdfs dfs -cat /mapreduce/task3output/part-r-00000
19/01/11 06:26:54 WARN util.NativeCodeLoader: Unable to load native-hadoop libr
Assam    0
Kerala   1
Madhya Pradesh  0
Maharashtra  0
Uttar Pradesh  3
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$

```